

Lab 7

Omar Youssif

Math 241, Week 9

```
# Put all necessary libraries here
library(tidyverse)
library(tidytext)
library(stringr)
library(wordcloud)
library(RColorBrewer)
# Ensure the textdata package is installed
if (!requireNamespace("textdata", quietly = TRUE)) {
  install.packages("textdata")
}
# Load the textdata package
library(textdata)

# Before knitting your document one last time, you will have to download the AFINN lexicon explicitly
lexicon_afinn()
```

```
## # A tibble: 2,477 x 2
##   word      value
##   <chr>    <dbl>
## 1 abandon    -2
## 2 abandoned  -2
## 3 abandons   -2
## 4 abducted   -2
## 5 abduction  -2
## 6 abductions -2
## 7 abhor      -3
## 8 abhorred   -3
## 9 abhorrent  -3
## 10 abhors    -3
## # i 2,467 more rows
```

```
lexicon_nrc()
```

```
## # A tibble: 13,872 x 2
##   word      sentiment
##   <chr>    <chr>
## 1 abacus   trust
## 2 abandon  fear
## 3 abandon  negative
## 4 abandon  sadness
## 5 abandoned anger
## 6 abandoned fear
## 7 abandoned negative
```

```
## 8 abandoned sadness
## 9 abandonment anger
## 10 abandonment fear
## # i 13,862 more rows
```

Due: Friday, March 29th at 5:30pm

Goals of this lab

1. Practice matching patterns with regular expressions.
2. Practice manipulating strings with `stringr`.
3. Practice tokenizing text with `tidytext`.
4. Practice looking at word frequencies.
5. Practice conducting sentiment analysis.

Problem 1: What's in a Name? (You'd Be Surprised!)

1. Load the `babynames` dataset, which contains yearly information on the frequency of baby names by sex and is provided by the US Social Security Administration. It includes all names with at least 5 uses per year per sex. In this problem, we are going to practice pattern matching!

```
library(babynames)
data("babynames")
#?babynames
```

- a. For 2000, find the ten most popular female baby names that start with the letter Z.

```
Fbaby2000 <- babynames %>% filter(year == 2000, sex == "F")
dfa <- top_n(Fbaby2000 %>% filter(Fbaby2000$name %in% str_subset(Fbaby2000$name, "^Z")),
            n = 10,
            wt = n)
dfa$dataset <- "dfa"
dfa
```

```
## # A tibble: 10 x 6
##   year sex  name      n      prop dataset
##   <dbl> <chr> <chr>   <int>   <dbl> <chr>
## 1  2000 F    Zoe     3785 0.00190 dfa
## 2  2000 F   Zoey    691 0.000346 dfa
## 3  2000 F   Zaria   568 0.000285 dfa
## 4  2000 F   Zoie    320 0.000160 dfa
## 5  2000 F  Zariah   168 0.0000842 dfa
## 6  2000 F   Zion    156 0.0000782 dfa
## 7  2000 F  Zainab   142 0.0000712 dfa
## 8  2000 F   Zara    121 0.0000607 dfa
## 9  2000 F  Zahra   113 0.0000566 dfa
## 10 2000 F  Zaira   103 0.0000516 dfa
```

- b. For 2000, find the ten most popular female baby names that contain the letter z.

```
dfb <- top_n(Fbaby2000 %>% filter(Fbaby2000$name %in% str_subset(Fbaby2000$name, "z")),
            n = 10,
            wt = n)
dfb$dataset <- "dfb"
dfb
```

```
## # A tibble: 10 x 6
##   year sex  name      n      prop dataset
```

```
##      <dbl> <chr> <chr>      <int>      <dbl> <chr>
## 1  2000 F      Elizabeth 15094 0.00757 dfb
## 2  2000 F      Mackenzie 6348 0.00318 dfb
## 3  2000 F      Mckenzie 2526 0.00127 dfb
## 4  2000 F      Makenzie 1613 0.000809 dfb
## 5  2000 F      Jazmin 1391 0.000697 dfb
## 6  2000 F      Jazmine 1353 0.000678 dfb
## 7  2000 F      Lizbeth 817 0.000410 dfb
## 8  2000 F      Eliza 759 0.000380 dfb
## 9  2000 F      Litzy 722 0.000362 dfb
## 10 2000 F      Esperanza 499 0.000250 dfb
```

c. For 2000, find the ten most popular female baby names that end in the letter z.

```
dfc <- top_n(Fbaby2000 %>% filter(Fbaby2000$name %in% str_subset(Fbaby2000$name, "z$")),
  n = 10,
  wt = n)
dfc$dataset <- "dfc"
dfc
```

```
## # A tibble: 11 x 6
##   year sex  name      n      prop dataset
##   <dbl> <chr> <chr>    <int>    <dbl> <chr>
## 1  2000 F    Luz      489 0.000245 dfc
## 2  2000 F   Beatriz 357 0.000179 dfc
## 3  2000 F  Mercedes 141 0.0000707 dfc
## 4  2000 F  Maricruz 96 0.0000481 dfc
## 5  2000 F    Liz      72 0.0000361 dfc
## 6  2000 F   Inez      69 0.0000346 dfc
## 7  2000 F  Odaliz    24 0.0000120 dfc
## 8  2000 F  Marycruz 23 0.0000115 dfc
## 9  2000 F    Cruz     19 0.00000952 dfc
## 10 2000 F   Deniz    16 0.00000802 dfc
## 11 2000 F   Taiz     16 0.00000802 dfc
```

d. Between your three tables in 1.a - 1.c, do any of the names show up on more than one list? If so, which ones? (Yes, I know you could do this visually but use some joins!)

Nope, we can also check this through inner joining all of them and doing all possible inner joins to catch all intersects, and we can see in the dataset “dfd” that we have 0 overlaps between all of them as it is empty.

```
dfd1 <- inner_join(dfa, dfb, by = "name")
dfd2 <- inner_join(dfa, dfc, by = "name")
dfd3 <- inner_join(dfb, dfc, by = "name")

dfd <- full_join(dfd1, dfd2, dfd3, by = "name")
dfd
```

```
## # A tibble: 0 x 21
## # i 21 variables: year.x.x <dbl>, sex.x.x <chr>, name <chr>, n.x.x <int>,
## #   prop.x.x <dbl>, dataset.x.x <chr>, year.y.x <dbl>, sex.y.x <chr>,
## #   n.y.x <int>, prop.y.x <dbl>, dataset.y.x <chr>, year.x.y <dbl>,
## #   sex.x.y <chr>, n.x.y <int>, prop.x.y <dbl>, dataset.x.y <chr>,
## #   year.y.y <dbl>, sex.y.y <chr>, n.y.y <int>, prop.y.y <dbl>,
## #   dataset.y.y <chr>
```

e. Verify that none of the baby names contain a numeric (0-9) in them.

This returns an empty set, so we know that none of the names have a numeric in them.

```
babynames %>% filter(babynames$name %in% str_subset(babynames$name, "[0-9]"))
```

```
## # A tibble: 0 x 5
```

```
## # i 5 variables: year <dbl>, sex <chr>, name <chr>, n <int>, prop <dbl>
```

- f. While none of the names contain 0-9, that doesn't mean they don't contain "one", "two", ..., or "nine". Create a table that provides the number of times a baby's name contained the word "zero", the word "one", ... the word "nine".

Notes:

- I recommend first converting all the names to lower case.
- If none of the baby's names contain the written number, there you can leave the number out of the table.
- Use `str_extract()`, not `str_extract_all()`. (We will ignore names where more than one of the words exists.)

Hint: You will have two steps that require pattern matching: 1. Subset your table to only include the rows with the desired words. 2. Add a column that contains the desired word.

```
lower_baby <- babynames
```

```
lower_baby$name <- lower_baby$name %>% str_replace_all(pattern = c("A" = "a", "B" = "b", "C" = "c", "D" = "d"))
```

```
num_baby <- lower_baby %>% filter(lower_baby$name %in% str_subset(lower_baby$name, "one|two|three|four|five|six|seven|eight|nine"))
```

```
dfd <- data.frame(
  one = sum(str_detect(num_baby$name, pattern = "one")),
  two = sum(str_detect(num_baby$name, pattern = "two")),
  three = sum(str_detect(num_baby$name, pattern = "three")),
  four = sum(str_detect(num_baby$name, pattern = "four")),
  five = sum(str_detect(num_baby$name, pattern = "five")),
  six = sum(str_detect(num_baby$name, pattern = "six")),
  seven = sum(str_detect(num_baby$name, pattern = "seven")),
  eight = sum(str_detect(num_baby$name, pattern = "eight")),
  nine = sum(str_detect(num_baby$name, pattern = "nine"))
)
```

```
print(dfd)
```

```
##      one two three four five six seven eight nine
## 1 10210 288   58    2    0 106   50   356  807
```

- g. Which written number or numbers don't show up in any of the baby names?

The written number "five" does not show up at all with "four" being very close to zero in terms of count too.

- h. Create a table that contains the names and their frequencies for the two least common written numbers.

```
num_baby %>% filter(str_detect(num_baby$name, pattern = "four"))
```

```
## # A tibble: 2 x 5
```

```
##   year sex  name      n      prop
##   <dbl> <chr> <chr>   <int>    <dbl>
## 1  1914 M    balfour    5 0.00000732
## 2  1928 M    balfour    5 0.00000438
```

```
num_baby %>% filter(str_detect(num_baby$name, pattern = "seven"))
```

```
## # A tibble: 50 x 5
```

```
##      year sex   name      n      prop
##      <dbl> <chr> <chr> <int>    <dbl>
## 1  1968 M     seven      5 0.00000282
## 2  1993 M     seven      7 0.00000339
## 3  1994 F     seven      5 0.00000257
## 4  1994 M     seven     10 0.00000491
## 5  1995 M     seven     11 0.00000547
## 6  1996 F     seven      7 0.00000365
## 7  1996 M     seven     10 0.00000499
## 8  1997 F     seven      7 0.00000367
## 9  1997 M     seven     22 0.0000110
## 10 1998 F     seven     14 0.00000722
## # i 40 more rows
```

- i. List out the names that contain no vowels (consider “y” to be a vowel).

```
lower_baby %>% filter(str_detect(lower_baby$name, pattern = "[^aeiouy]+$"))
```

```
## # A tibble: 1,260 x 5
##      year sex   name      n      prop
##      <dbl> <chr> <chr> <int>    <dbl>
## 1  1880 M     wm      14 0.000118
## 2  1881 M     wm      15 0.000139
## 3  1882 M     wm      18 0.000148
## 4  1883 M     wm      16 0.000142
## 5  1884 M     wm      22 0.000179
## 6  1885 M     wm      23 0.000198
## 7  1885 M     ng       5 0.0000431
## 8  1886 M     wm      15 0.000126
## 9  1887 M     wm      15 0.000137
## 10 1888 M     wm      25 0.000192
## # i 1,250 more rows
```

(I’m not sure what dataset to use here, the question is a little unclear to me. Can’t tell whether all baby names or the written number one. I checked and the written number one yields 0 names).

Problem 2: Tidying the “Call of the Wild”

Did you read “Call of the Wild” by Jack London? If not, [read the first paragraph of its wiki page](#) for a quick summary and then let’s do some text analysis on this classic! The following code will pull the book into R using the `gutenbergr` package.


```
library(gutenbergr)
wild <- gutenberg_download(215)
```

- a. Create a tidy text dataset where you tokenize by words.

```
p2dfa <- wild %>% unnest_tokens(output = word, input = text) %>% select(-gutenberg_id)
```

- b. Find the frequency of the 20 most common words. First, remove stop words.

```
p2dfa %>%
  anti_join(stop_words) %>%
  count(word, sort = TRUE) %>%
  filter(n > 39) %>%
  ggplot(aes(y = fct_reorder(word, n), x = n, fill = n)) +
  geom_col() +
  guides(fill = FALSE)
```



```
p2df <- p2dfa %>% anti_join(stop_words) %>% count(word, sort = TRUE) %>% filter(n > 39)
```

c. Create a bar graph and a word cloud of the frequencies of the 20 most common words.

```
pal <- brewer.pal(9, "Set1")

wordcloud(p2df$word,
  scale = c(3, 1),
  rot.per = .5, colors = pal,
  min.freq = 0, random.order = FALSE)
```

eyes buck's
spitz
headtime days
trail thornton
john lifebuckday
camp sleds dogfeet
half
françois night

- d. Explore the sentiment of the text using three of the sentiment lexicons in `tidytext`. What does your analysis say about the sentiment of the text?

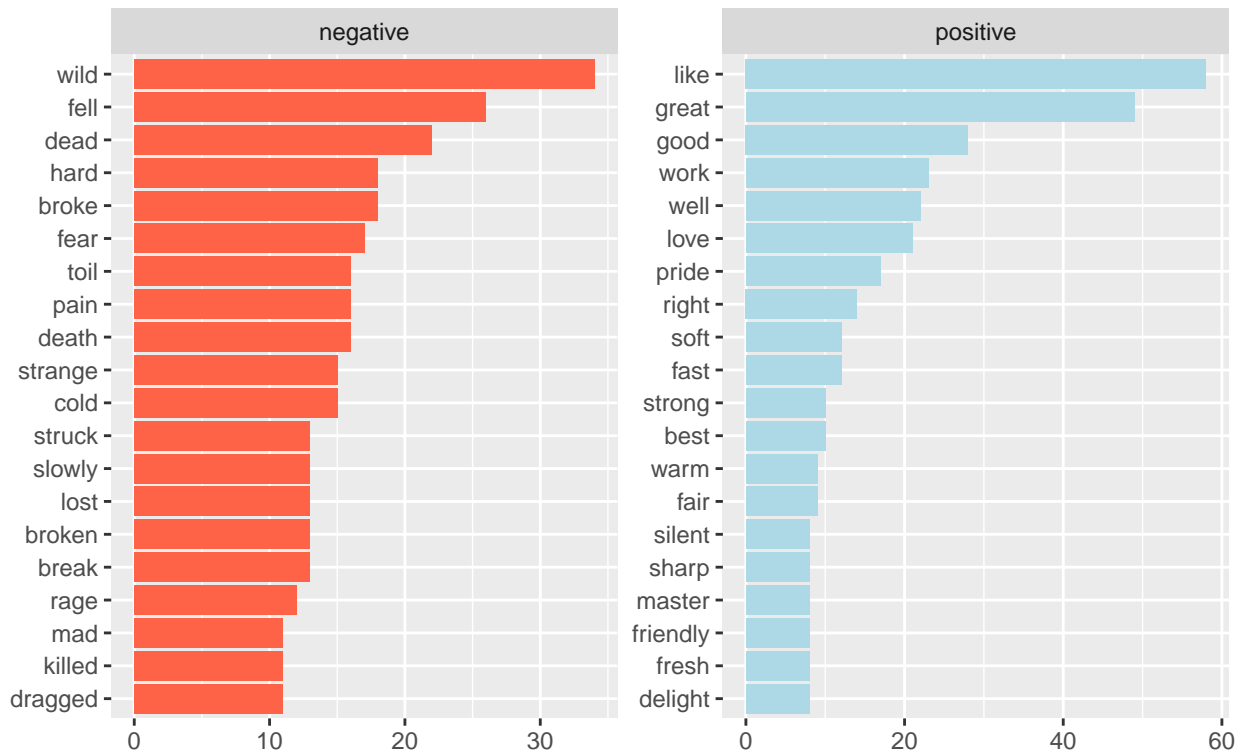
Notes:

- Make sure to NOT remove stop words this time.
- `afinn` is a numeric score and should be handled differently than the categorical scores.

```
p2dfa %>%
  inner_join(get_sentiments("bing"), by = "word") %>%
  count(sentiment, word, sort = TRUE) %>%
  group_by(sentiment) %>%
  slice_head(n = 20) %>%
  ggplot(aes(y = fct_reorder(word, n), x = n, fill = sentiment)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~sentiment, scales = "free") +
  labs(
    title = "Sentiment and frequency of words in Call of the Wild by Jack London?",
    subtitle = "Bing lexicon",
    y = NULL, x = NULL
  ) +
  scale_fill_manual(values = c("tomato", "lightblue"))
```

Sentiment and frequency of words in Call of the Wild by Jack London?

Bing lexicon



```
p2dfa %>%
inner_join(get_sentiments("nrc"), by = "word") %>%
mutate(
  sentiment = fct_relevel(
    sentiment, "positive", "anticipation", "joy", "surprise", "trust",
    "negative", "anger", "disgust", "fear", "sadness"
  ),
  sentiment_binary = if_else(sentiment %in% c("positive", "anticipation", "joy", "surprise", "trust")
) %>%
count(sentiment_binary, sentiment, word, sort = TRUE) %>%
group_by(sentiment) %>%
slice_head(n = 10) %>%
ggplot(aes(y = fct_reorder(word, n), x = n, fill = sentiment_binary)) +
geom_col() +
guides(fill = FALSE) +
facet_wrap(~sentiment, scales = "free_y", ncol = 5) +
labs(
  title = "Sentiment and frequency of words in Call of the Wild by Jack London?",
  subtitle = "NRC lexicon",
  y = NULL, x = NULL
) +
scale_x_continuous(breaks = c(0, 200)) +
theme_minimal(base_size = 11) +
scale_fill_manual(values = c("tomato", "lightblue"))
```


Sentiment and frequency of words in Call of the Wild by Jack London? NRC lexicon




```
p2dfe <- p2dfa %>% inner_join(get_sentiments("afinn")) %>%
  group_by(word) %>%
  summarise(sentiment = sum(value))

ggplot(data = p2dfe, mapping = aes(x = sentiment)) +
  geom_bar() +
  labs(
    title = "Sentiment and frequency of words in Call of the Wild by Jack London?",
    subtitle = "AFINN lexicon")
```

Sentiment and frequency of words in Call of the Wild by Jack London?

AFINN lexicon



The text seems to lean a little more towards the negative side, particularly when looking at the positive words and some possibly being connector words, not actual positive words.

- e. If you didn't do so in 2.d, compute the average sentiment score of the text using `afinn`. Which positive words had the biggest impact? Which negative words had the biggest impact?

```
p2dfe %>% summarize(avg = mean(sentiment))
```

```
## # A tibble: 1 x 1
##   avg
##   <dbl>
## 1 -1.12
```

```
p2dfe %>% arrange(desc(sentiment)) %>% slice_head(n = 5)
```

```
## # A tibble: 5 x 2
##   word      sentiment
##   <chr>      <dbl>
## 1 great      147
## 2 like       116
## 3 good        84
## 4 love        63
## 5 strength    42
```

```
p2dfe %>% arrange(sentiment) %>% slice_head(n = 5)
```

```
## # A tibble: 5 x 2
##   word      sentiment
##   <chr>      <dbl>
## 1 no        -95
```

```
## 2 dead      -66
## 3 fire      -66
## 4 cried     -42
## 5 lost      -39
```

Positive: great, like, good, love, strength Negative: no, dead, fire, cried, lost

- f. You should have found that “no” was an important negative word in the sentiment score. To know if that really makes sense, let’s turn to the raw lines of text for context. Pull out all of the lines that have the word “no” in them. Make sure to not pull out extraneous lines (e.g., a line with the word “now”).

```
dff <- wild %>%
  unnest_tokens(output = words, input = text, token = "lines") %>%
  select(-gutenberg_id) %>%
  mutate(id = row_number())

dff %>% group_by(id) %>%
  filter(str_count(words, "\\s") > 1) %>%
  filter(str_detect(words, "\\bno\\b"))
```

```
## # A tibble: 94 x 2
## # Groups:   id [94]
##   words                                     id
##   <chr>                                     <int>
## 1 manuel’s treachery. no one saw him and buck go off through the orchard      86
## 2 solitary man, no one saw them arrive at the little flag station known      88
## 3 that it was the club, but his madness knew no caution. a dozen times he    217
## 4 "he’s no slouch at dog-breakin’, that’s wot i say," one of the men on      233
## 5 all, that he stood no chance against a man with a club. he had learned      252
## 6 in the red sweater. "and seem’ it’s government money, you ain’t got no      280
## 7 animal. the canadian government would be no loser, nor would its            284
## 8 while he developed no affection for them, he none the less grew            297
## 9 the other dog made no advances, nor received any; also, he did not         311
## 10 heart of civilization and flung into the heart of things primordial. no    338
## # i 84 more rows
```

- g. Draw some conclusions about how “no” is used in the text.

It seems to be used mostly as a synonym of “regardless of” but always in a negative connotation, so I see why the lexicons thought it’s a very negative word. Though this is a overgeneralization and there are some cases where that’s not the case, but it can’t separate them, so it’s incredibly boosted.

- h. We can also look at how the sentiment of the text changes as the text progresses. Below, I have added two columns to the original dataset. Now I want you to do the following wrangling:

- Tidy the data (but don’t drop stop words).
- Add the word sentiments using `bing`.
- Count the frequency of sentiments by index.
- Reshape the data to be wide with the count of the negative sentiments in one column and the positive in another, along with a column for index.
- Compute a sentiment column by subtracting the negative score from the positive.

```
wild_time <- wild %>%
  mutate(line = row_number(), index = floor(line/45) + 1)
```

```
dfh1 <- wild_time %>%
  unnest_tokens(output = word, input = text) %>%
  select(-gutenberg_id) %>%
```

```

inner_join(get_sentiments("bing"), by = "word") %>%
count(sentiment, index, sort = TRUE)

dfh2 <- wild_time %>%
  unnest_tokens(output = word, input = text) %>%
  select(-gutenberg_id) %>%
  inner_join(get_sentiments("bing"), by = "word") %>%
  count(sentiment, index, sort = TRUE) %>%
  pivot_wider(names_from = sentiment, values_from = n) %>%
  mutate(vibe = positive - negative)

```

dfh1

```

## # A tibble: 138 x 3
##   sentiment index      n
##   <chr>      <dbl> <int>
## 1 negative     17     38
## 2 negative     19     35
## 3 negative     22     35
## 4 negative     63     34
## 5 negative     44     33
## 6 negative     21     32
## 7 negative     18     30
## 8 negative     33     30
## 9 negative     40     30
## 10 negative    42     30
## # i 128 more rows

```

dfh2

```

## # A tibble: 69 x 4
##   index negative positive  vibe
##   <dbl>      <int>      <int> <int>
## 1     17         38         11  -27
## 2     19         35          5  -30
## 3     22         35         16  -19
## 4     63         34         13  -21
## 5     44         33          7  -26
## 6     21         32         17  -15
## 7     18         30          6  -24
## 8     33         30          9  -21
## 9     40         30         14  -16
## 10    42         30         11  -19
## # i 59 more rows


```

i. Create a plot of the sentiment scores as the text progresses.

```

ggplot(data = dfh2, mapping = aes(x = index, y = vibe)) +
  geom_line()

```




- j. The choice of 45 lines per chunk was pretty arbitrary. Try modifying the index value a few times and recreating the plot in i. Based on your plots, what can you conclude about the sentiment of the novel as it progresses?

```
wild_time2 <- wild %>%
  mutate(line = row_number(), index = floor(line/15) + 1)

dfj <- wild_time2 %>%
  unnest_tokens(output = word, input = text) %>%
  select(-gutenberg_id) %>%
  inner_join(get_sentiments("bing"), by = "word") %>%
  count(sentiment, index, sort = TRUE) %>%
  pivot_wider(names_from = sentiment, values_from = n) %>%
  mutate(vibe = positive - negative)

ggplot(data = dfj, mapping = aes(x = index, y = vibe)) +
  geom_line()
```



It seems generally stagnant (trend-wise) though. it doesn't look as extremely negative as the other line makes it out to be. Though it seems like it goes up and down significantly in short periods of times, which makes a large index division lean negative.

k. Let's look at the bigrams (2 consecutive words). Tokenize the text by bigrams.

```
dfk <- wild %>%
  unnest_tokens(output = bigrams, input = text, token = "ngrams", n = 2) %>%
  select(-gutenberg_id) %>%
  mutate(id = row_number()) %>%
  drop_na()
```

dfk

```
## # A tibble: 29,442 x 2
##   bigrams      id
##   <chr>      <int>
## 1 the call        6
## 2 call of         7
## 3 of the          8
## 4 the wild        9
## 5 by jack       11
## 6 jack london   12
## 7 chapter i     19
## 8 i into        20
## 9 into the      21
## 10 the primitive 22
## # i 29,432 more rows
```

1. Produce a sorted table that counts the frequency of each bigram and notice that stop words are still an issue.

```
dfk %>%  
  unnest_tokens(word, bigrams, drop = FALSE) %>%  
  anti_join(stop_words) %>%  
  group_by(bigrams) %>%  
  summarize(n = n()) %>%  
  arrange(desc(n))
```

```
## # A tibble: 15,351 x 2  
##   bigrams      n  
##   <chr>    <int>  
## 1 john thornton    60  
## 2 sol leks        56  
## 3 the sled        44  
## 4 the dogs        35  
## 5 buck was        28  
## 6 his feet        26  
## 7 and buck        23  
## 8 half breed      22  
## 9 red sweater     22  
## 10 the team       21  
## # i 15,341 more rows
```